

Caleb Ziems

Curriculum Vitae

✉ cziems@stanford.edu

🌐 www.calebziems.com

🐦 [cjziems](#)

🌐 [cjziems](#)

*I work on **socially competent** and **equitable NLP** systems that are more aware of **social norms** and **robust to language variation**. My work also supports **safety** and **wellbeing** in online communities.*

Education

- 2023–Present **Stanford University**, *Ph.D. in Computer Science*.
NLP Group
Advisor: Diyi Yang
- 2020–2023 **Georgia Institute of Technology**, *M.S. in Computer Science*.
School of Interactive Computing
Advisor: Diyi Yang
Cumulative GPA: 3.95/4.0
- 2016–2020 **Emory University**, *B.S. in Computer Science (highest honors), B.A. in Mathematics*.
Thesis: *Towards More Robust Methods of Cyberbullying Detection*
Advisor: Ymir Vigfusson
Cumulative GPA: 4.0/4.0

Publications

- 2023 **Ziems, C.**, Held, W., Shaikh, O., Chen, J., Zhang, Z., & Yang, D. (2023). Can Large Language Models Transform Computational Social Science? In *Computational Linguistics*, 49(4). [PDF](#)
- 2023 Kruk, J., **Ziems, C.**, & Yang, D. (2023). Impressions: Visual Semiotics and Aesthetic Impact Understanding. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. [PDF](#)
- 2023 Li, E., Shi, T., **Ziems, C.**, Kan, M., Chen, N., Liu, Z., & Yang, D. (2023). CoAnnotating: Uncertainty-Guided Work Allocation between Human and Large Language Models for Data Annotation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. [PDF](#)
- 2023 **Ziems, C.***, Held, W.*, Yang, J., & Yang, D. (2023). Multi-VALUE: A Framework for Cross-Dialectal English NLP. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*. [PDF](#)
- 2023 **Ziems, C.**, Dwivedi-Yu, J., Wang, Y., Halevy, A. Y., & Yang, D. (2023). NormBank: A Knowledge Bank of Situational Social Norms. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*. [PDF](#)
- 2023 Shaikh, O.*, **Ziems, C.***, Held, W., Pariani, A. J., Morstatter, F., & Yang, D. (2023). Modeling Cross-Cultural Pragmatic Inference with Codenames Duet. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics: Findings*. [PDF](#)

- 2023 **ACL** Held, W., **Ziems, C.**, & Yang, D. (2023). TADA: Task-Agnostic Dialect Adapters for English. In Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics: Findings. [PDF](#)
- 2022 **ACL** **Ziems, C.**, Chen, J., Harris, C., Anderson, A., & Yang, D. (2022). VALUE: Understanding Dialect Disparity in NLU. In Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing. [PDF](#)
- 2022 **ACL** **Ziems, C.**, Yu, J. A., Wang, Y., Halevy, A. Y., & Yang, D. (2022). The Moral Integrity Corpus: A Benchmark for Ethical Dialogue Systems. In Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing. [PDF](#)
- 2022 **ACL** **Ziems, C.***, Li, M.*, Zhang, A., & Yang, D. (2022). Inducing Positive Perspectives with Text Reframing. In Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing. [PDF](#)
[Outstanding Paper Award.](#)
- 2021 **EMNLP** **Ziems, C.***, ElSherief, M.*, Muchlinski, D., Anupindi, V., Seybolt, J., De Choudhury, M., & Yang, D. (2021). Latent Hatred: A Benchmark for Understanding Implicit Hate Speech. In Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing. [PDF](#)
- 2021 **EMNLP Findings** **Ziems, C.**, Yang, D. (2021). To Protect and To Serve? Analyzing Entity-Centric Framing of Police Violence. In Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing: Findings. [PDF](#)
- 2021 **ASONAM** He, B., **Ziems, C.**, Soni, S., Ramakrishnan, N., Yang, D., & Kumar, S. (2021). Racism is a Virus: Anti-Asian Hate and Counterhate in Social Media during the COVID-19 Crisis. In 2021 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining. IEEE. [PDF](#)
- 2020 **ICWSM** **Ziems, C.**, Vigfusson, Y., & Morstatter, F. (2020) Aggressive, Repetitive, Intentional, Visible, and Imbalanced: Refining Representations for Cyberbullying Classification. In Proceedings of the 14th International AAAI Conference on Web and Social Media. [PDF](#).
[Honorable Mention for Best Paper Award.](#)

Experience

- May 2022 - **Meta AI Integrity**, *Research Intern.*
- August 2022 Manager: Alon Halevy
Peers: Jane Yu, Yi-Chia Wang
- May 2021 - **Meta AI Integrity**, *Research Intern.*
- Sep 2021 Manager: Alon Halevy
Peers: Jane Yu, Yi-Chia Wang
- August 2020 - **Georgia Tech SALT Lab**, *Graduate Research Assistant.*
- Present Advisor: Diyi Yang
- March 2020 - **Georgia Tech CLAWS Research Group**, *Undergraduate Research Assistant.*
- May 2020 Advisor: Srijan Kumar

- June 2019 - **USC/ISI MINDS Research Group**, *NSF REU Fellow*.
August 2019 Advisor: Fred Morstatter
June 2018 - **Stanford InfoLab**, *NSF REU Fellow*.
August 2018 Advisor: Srijan Kumar
PI: Jure Leskovec
August 2017 - **Emory Concept Mining Lab**, *Undergraduate Research Assistant*.
May 2018 Advisor: Phillip Wolff

Teaching

- January 2022 **CS 6471 Computational Social Science**, *Graduate Teaching Assistant*.
- May 2022
 - Co-designed the CS 6471 curriculum
 - Co-designed the course schedule of assigned readings
 - Co-designed project report and presentation rubrics
 - Designed quizzes and bi-weekly reading response questions
 - Graded reports and presentations
- August 2017 - **Emory Office for Undergraduate Education**, *EPASS Peer Tutor*.
May 2018
 - Worked one-on-one with 10 to 20 different undergraduates each semester
 - Explained Calculus I and II concepts, designed problem sets, and offered individual feedback

Mentoring

- Fall 2023 - **Vyoma Raman**, *Stanford University (M.S.)*
Present *Topic*: Power Dynamics in NLP Benchmarking
- Fall 2023 - **Seong Hee Lee**, *Stanford University (M.S.)*
Present *Topic*: LLM Social Skill Learning
- Summer 2023 **Chijioke Mgbahurike**, *Stanford University (B.S.)*
- Present *Topic*: Dialectal Bias in ASR Systems
Result: One presentation at the INSPIRE Research Poster Session.
- Fall 2022 - **Julia Kruk**, *Georgia Institute of Technology (M.S.)*
Summer 2023 *Topic*: Computational Aesthetics and Photojournalism
Results: One EMNLP 2023 Paper.
- Spring 2022 - **Aryan Pariani**, *Georgia Institute of Technology (B.S.)*
Spring 2023 *Topic*: Pragmatic Reasoning in Language Games
Results: One ACL 2023 Findings Publication.
- Spring 2021 - **Minzhi (Ella) Li**, *National University of Singapore (B.S., Ph.D)*
Summer 2023 *Topic*: Positive Reframing; Human-LLM Co-Annotation
Results: One ACL 2022 **Outstanding Paper Award**; One EMNLP 2023 Paper.
- Fall 2021 - **Taiwei Shi**, *Georgia Institute of Technology (B.S.)*
Spring 2022; *Topics*: Dialect-Aware NLP; Human-LLM Co-Annotation
Summer 2023 *Results*: Undergraduate Thesis; One EMNLP 2023 Paper.

Awards and Honors

- Dec 2022 **Honorable Mention for Best Project**, *CS 7643 Deep Learning @ Georgia Tech*.
May 2022 **Outstanding Paper**, *ACL 2022*.
March 2021 **NSF Graduate Research Fellowship**.

- June 2020 **Honorable Mention for Best Paper**, *ICWSM 2020*.
- May 2020 **Summa Cum Laude**, *Emory University Computer Science Department*.
- May 2020 **Academic Excellence Award**, *Emory University Computer Science Department*.
- November 2018 **Phi Beta Kappa**.
Elected as a Junior
- August 2017 **Norman C. and Henry J. Miller Dean's Achievement Scholarship**.
\$30,000 scholarship recognizing academic achievement, leadership, and community involvement

Skills

- Crowdsourcing Amazon Mechanical Turk, Prolific, Upwork
- Languages Python, MATLAB, R, Java, C, Prolog, Scheme (LISP), JavaScript
- Libraries PyTorch, scikit-learn, NumPy, Pandas, spaCy, NetworkX, BeautifulSoup
- Software Bash, Vim, Gephi, Jupyter, Eclipse
- Tools HTML/CSS, PostgreSQL, Github, \LaTeX , Slurm

Service

- Reviewing ACM CSCW (2021);
ACL / ARR Review (2021, 2022, 2023)
- Volunteer EMNLP (2021); ACL (2023)

Talks and Miscellanea

- October 9, 2023 **Speaker** at the AWS Responsible AI Group.
Talk: Multi-VALUE: A Framework for Cross-Dialectal English NLP
- June 19, 2023 **Speaker** at the Summer Institutes in Computational Social Science (SICSS).
Talk: Can Large Language Models Transform Computational Social Science?
- October 4, 2022 **Panel Member** for CS 7001 Introduction to Graduate Studies at Georgia Tech.
Topic: National Fellowships
- February 24, 2022 **Student Host** for Yian Yin, Faculty Candidate in the School of Interactive Computing
Topic: Science of Science and Network Science
- February 14, 2022 **Speaker** at IST-Unbabel Seminar
Talk: Social NLP for Positive and Civil Interactions
- November 17, 2021 **Guest Lecturer** for CS 4650 Natural Language Processing at Georgia Tech
Topic: Dialect-Aware NLP + Computational Social Science
- October 5, 2021 **Panel Member** for CS 7001 Introduction to Graduate Studies at Georgia Tech.
Topic: National Fellowships